


Processes

Day3: Sep 2021

Kiran Waghmare






Agenda: Processes

- ❑ Preemptive and non preemptive
- ❑ Process mgmt
- ❑ Process life cycle
- ❑ Schedulers
- ❑ Scheduling algorithms
- ❑ Creation of fork, waitpid, exec system calls
- ❑ Orphan and zombie





Process in Memory





Process Concept

- ❑ An operating system executes a variety of programs:
 - ❑ Batch system – jobs
 - ❑ Time-shared systems – user programs or tasks
- ❑ Textbook uses the terms *job* and *process* almost interchangeably
- ❑ **Process – a program in execution; process execution must progress in sequential fashion**
- ❑ A process includes:
 - ❑ program counter
 - ❑ stack
 - ❑ data section






Process State

- ❑ As a process executes, it changes *state*
 - ❑ **new**: The process is being created
 - ❑ **running**: Instructions are being executed
 - ❑ **waiting**: The process is waiting for some event to occur
 - ❑ **ready**: The process is waiting to be assigned to a processor
 - ❑ **terminated**: The process has finished execution





Diagram of Process State





Process Control Block (PCB)

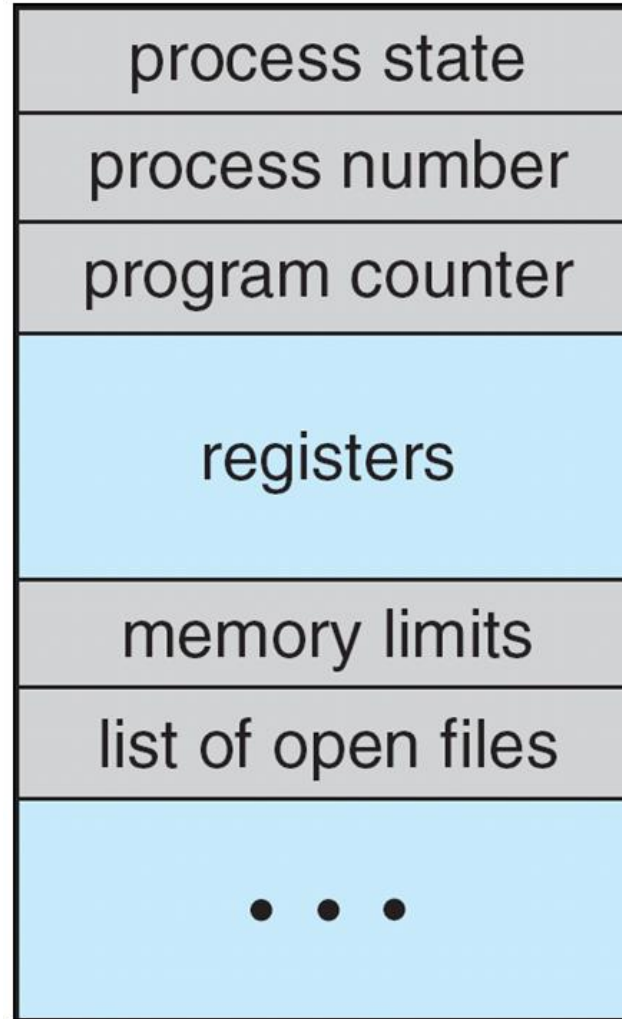
Information associated with each process

- ❑ Process state
- ❑ Program counter
- ❑ CPU registers
- ❑ CPU scheduling information
- ❑ Memory-management information
- ❑ Accounting information
- ❑ I/O status information





Process Control Block (PCB)





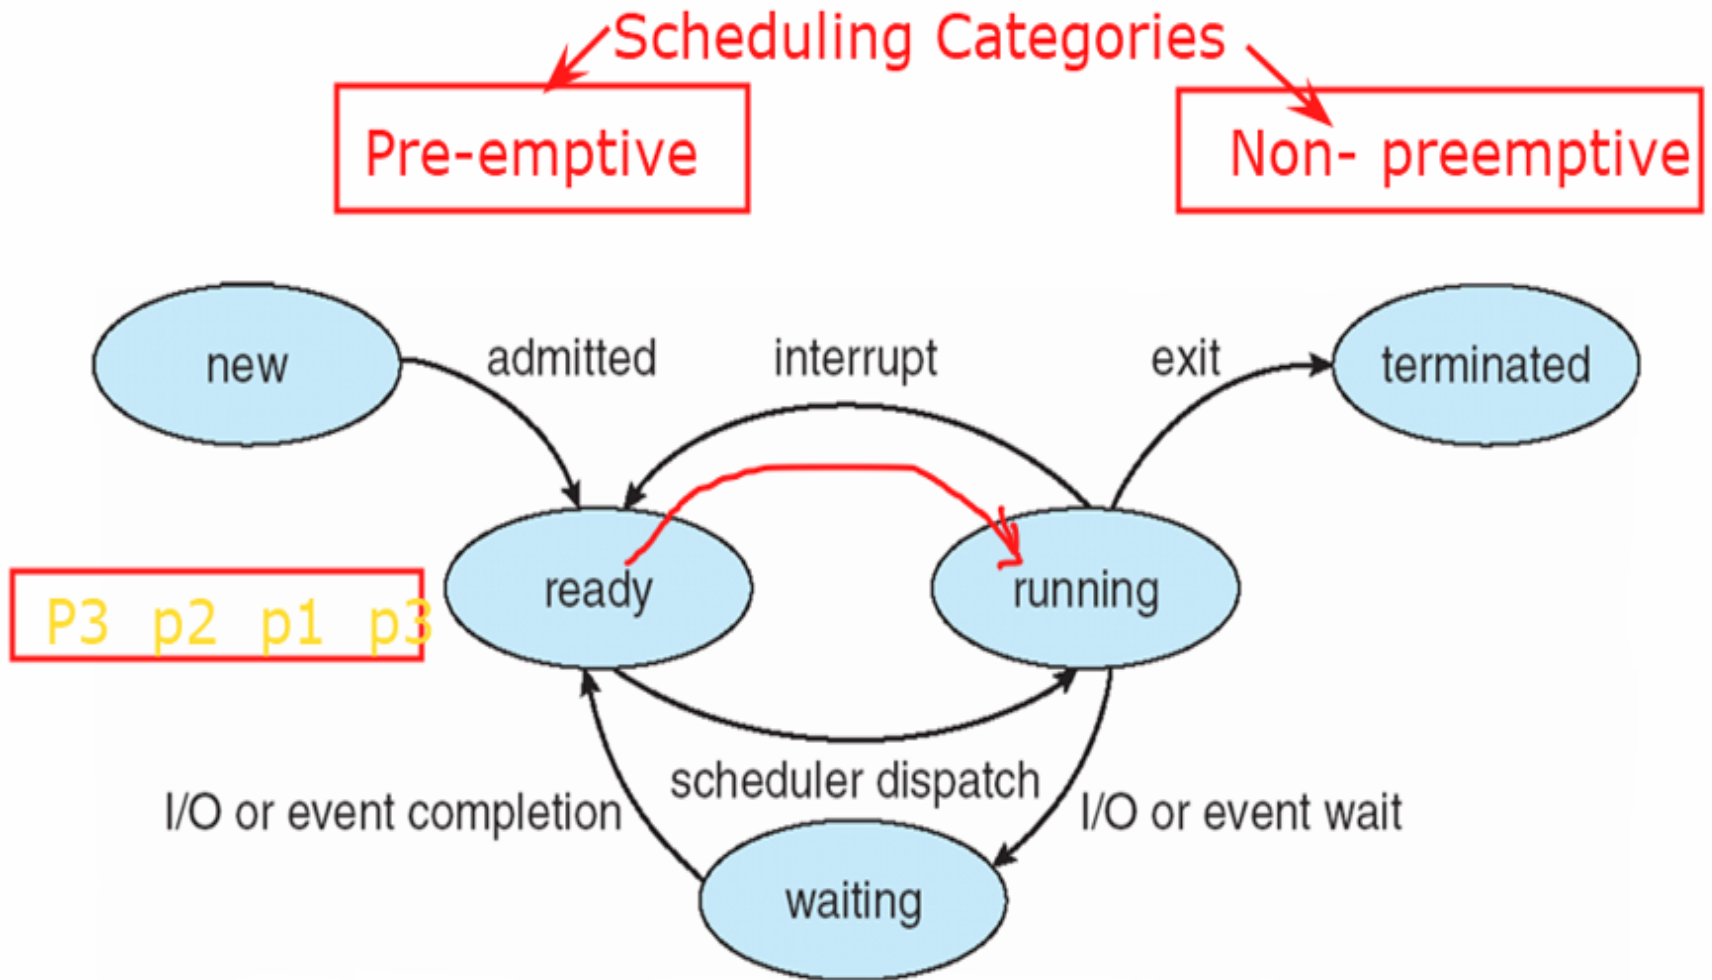
Process Scheduling

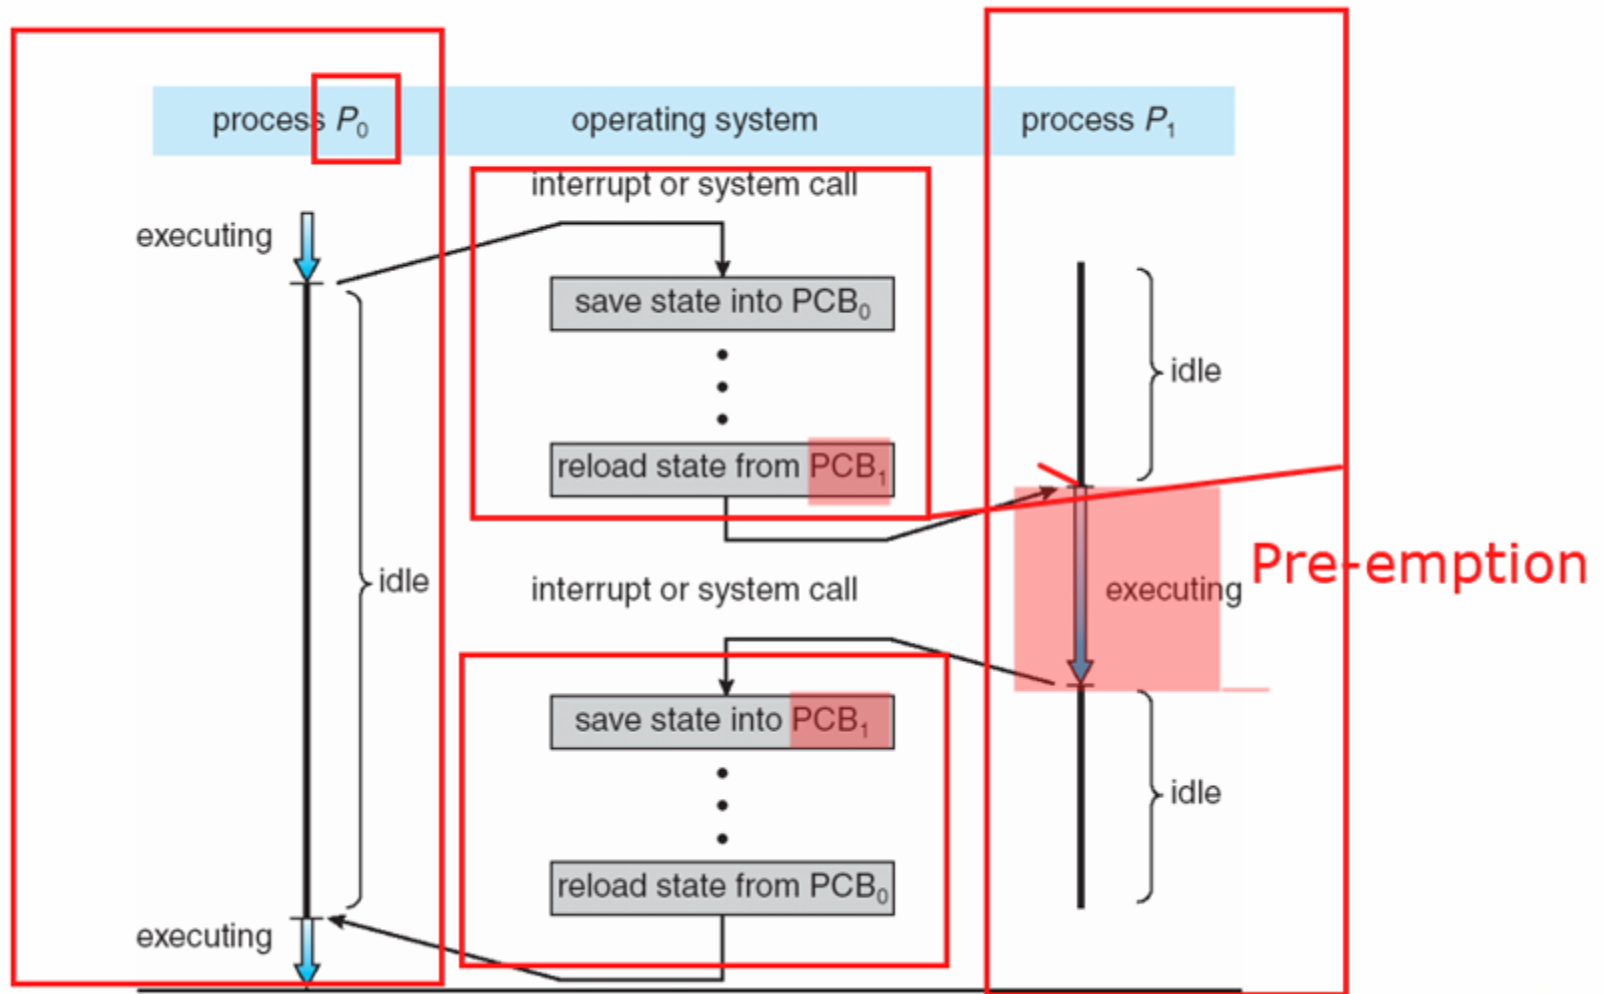
- When there are **two or more runnable processes** then it is decided by the Operating system which one to run first then it is referred to as Process Scheduling.
- A scheduler is **used to make decisions** by using some scheduling algorithm.
- Given below are the properties of a **Good Scheduling Algorithm**:
 - **Response time** should be **minimum** for the users.
 - The **number of jobs processed per hour should be maximum** i.e Good scheduling algorithm should give maximum throughput.
 - The **utilization of the CPU should be 100%**.
 - Each process should get a **fair share of the CPU**.

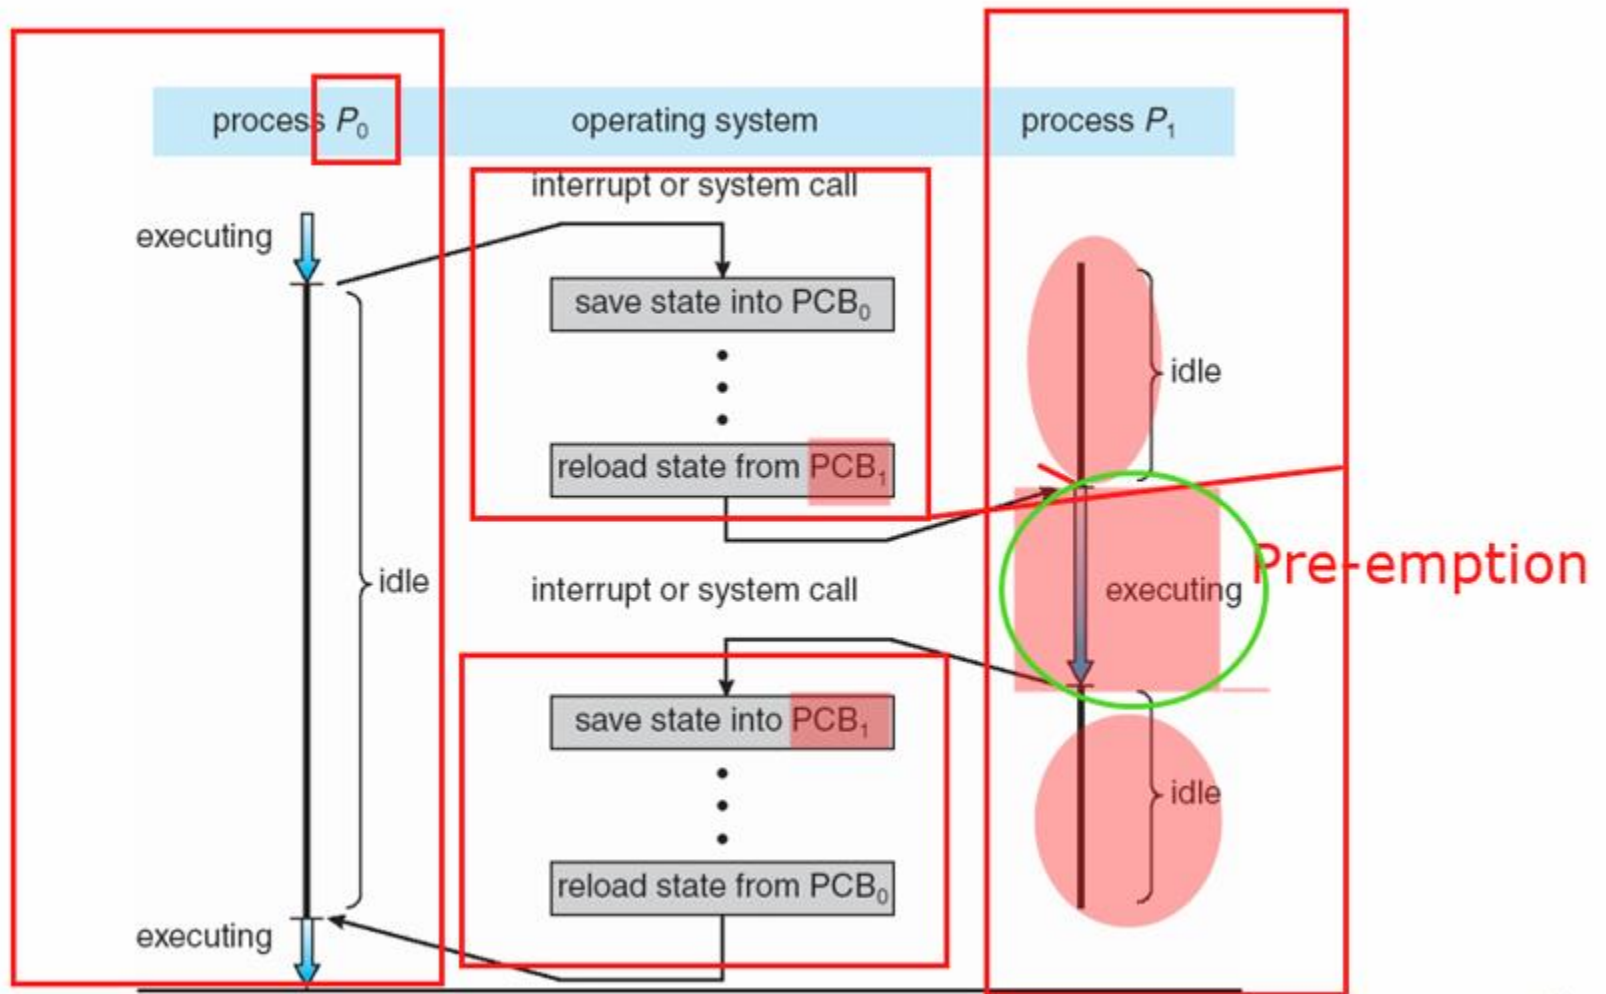




Diagram of Process State










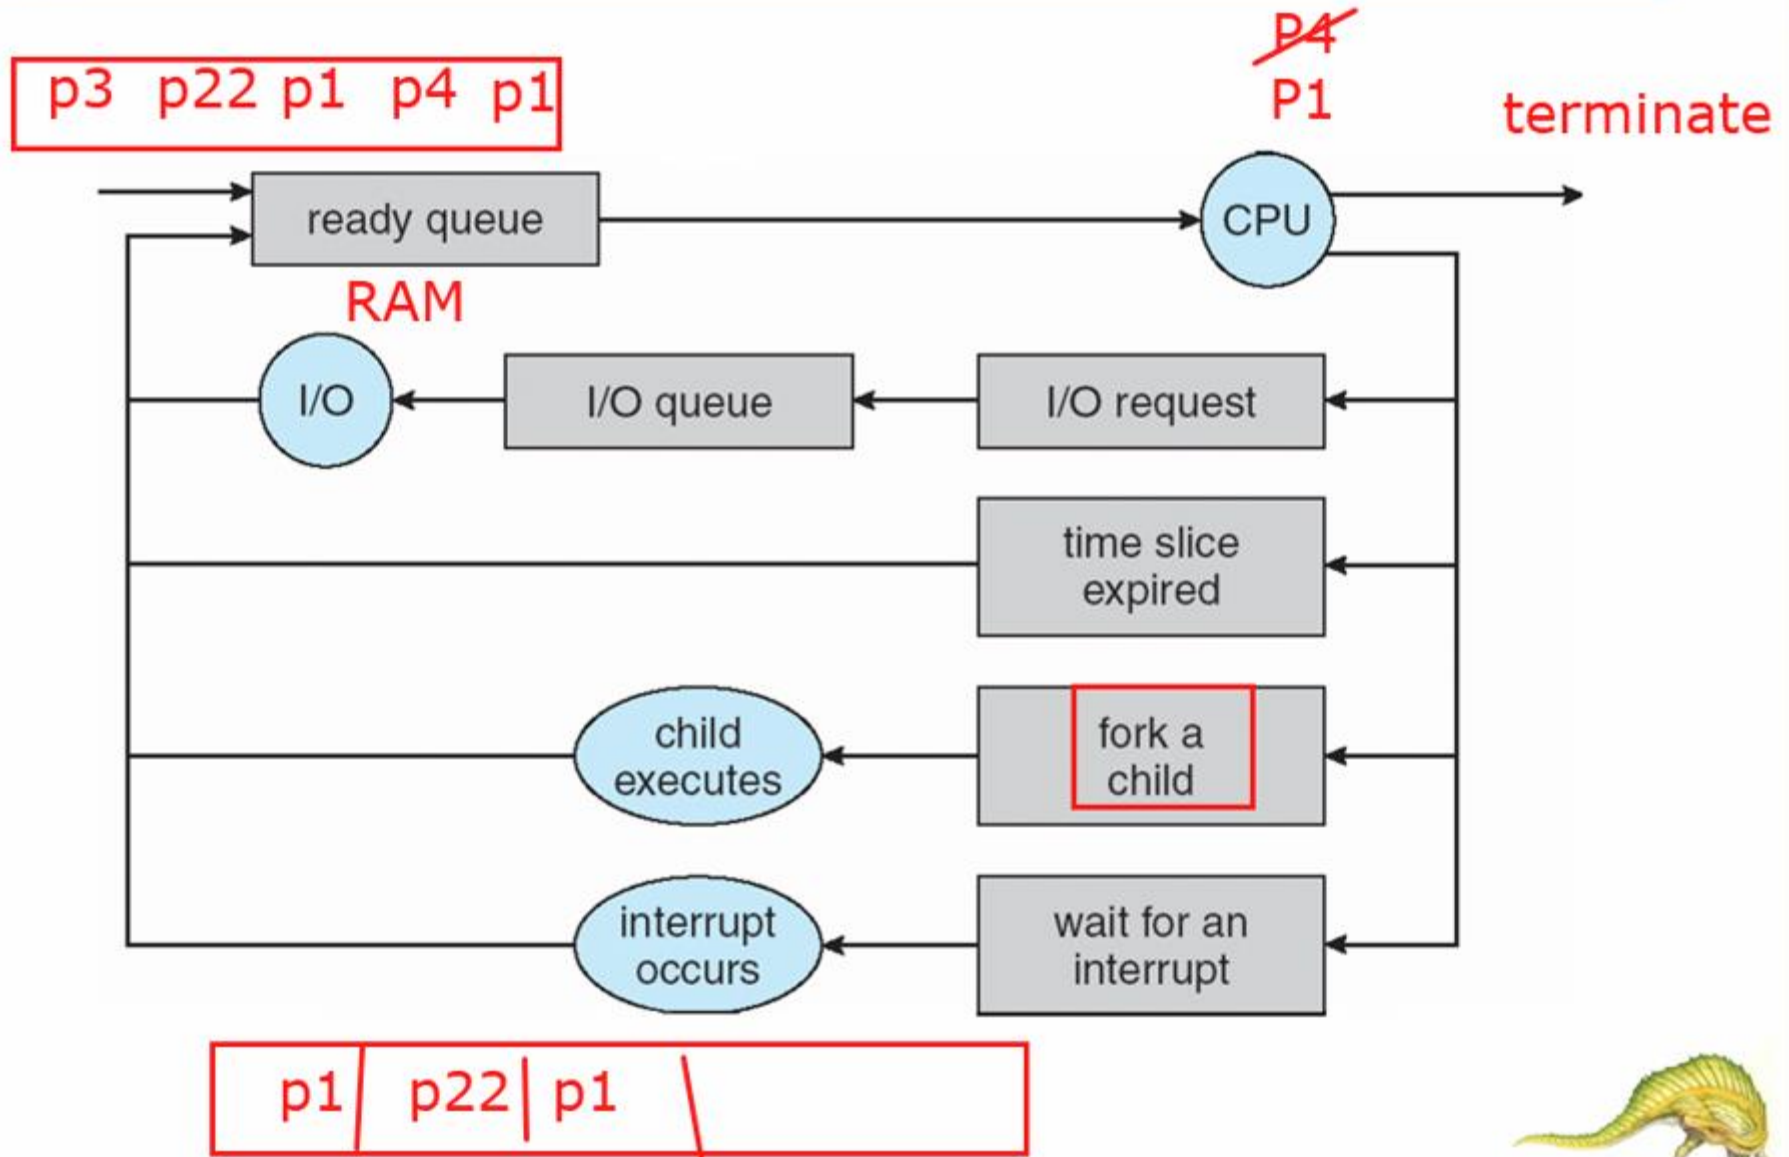
Process Scheduling Queues

- ❑ **Job queue** – set of all processes in the system
- ❑ **Ready queue** – set of all processes **residing in main memory**, ready and waiting to execute
- ❑ **Device queues** – set of processes **waiting for an I/O** device
- ❑ Processes migrate among the various queues



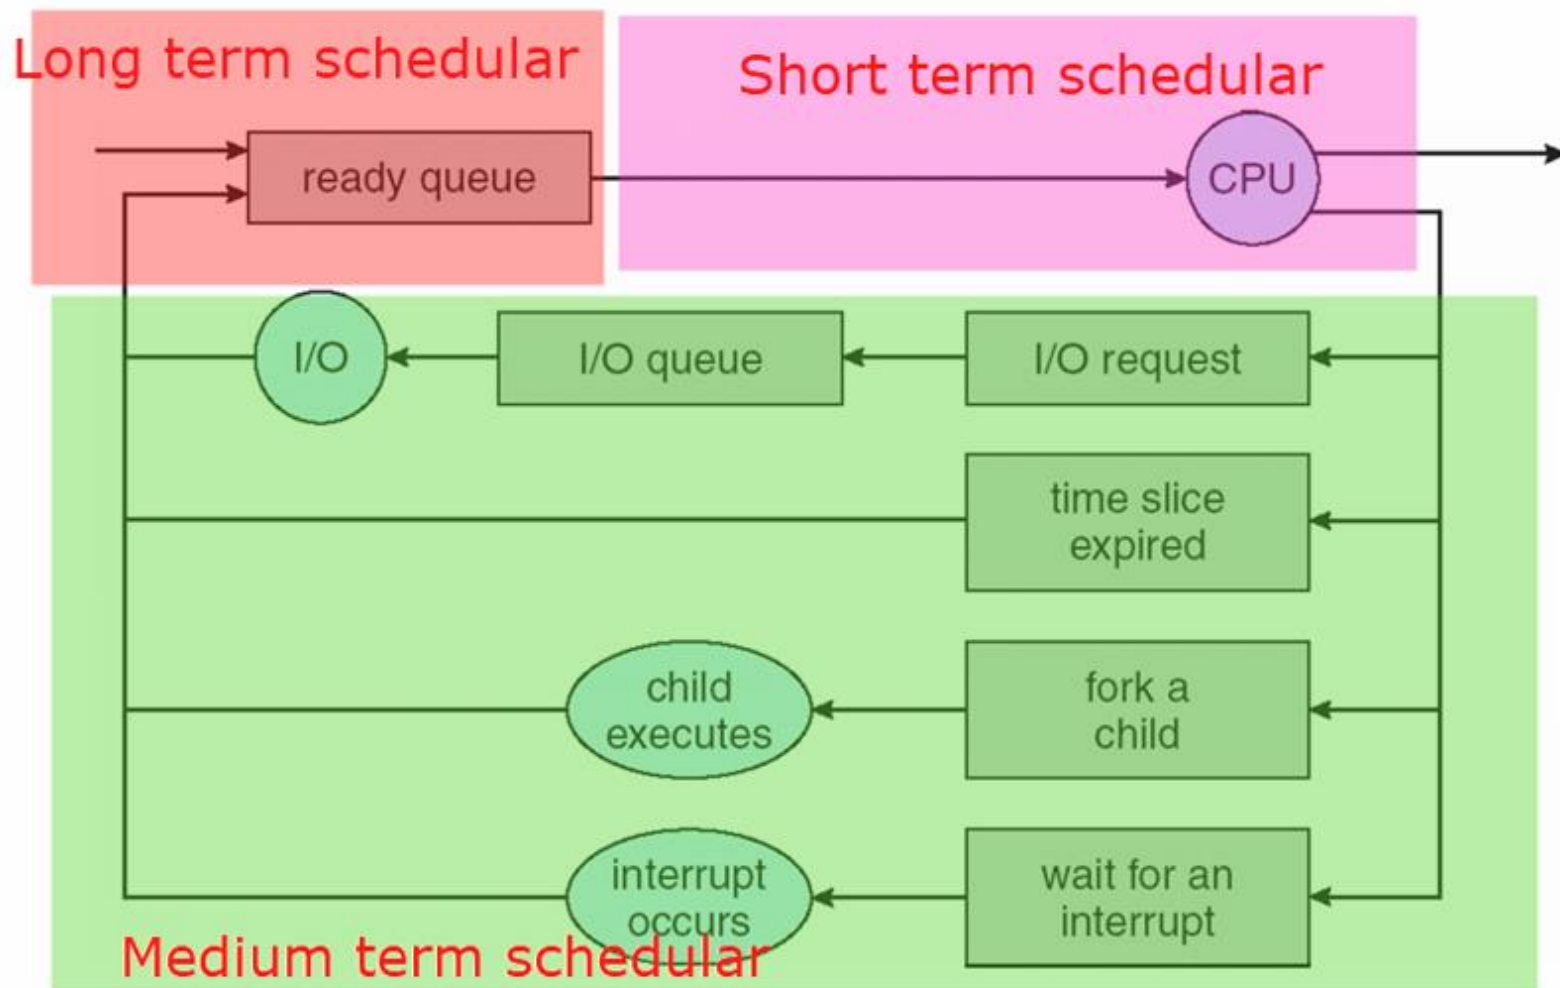


Representation of Process Scheduling





Representation of Process Scheduling






Types of Schedulers

- There are three types of schedulers available:
- Long Term Scheduler
- Short Term Scheduler
- Medium Term Scheduler





Addition of Medium Term Scheduling





Context Switch

- When CPU switches to another process, the system must **save the state of the old process and load the saved state for the new process** via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
- Time dependent on hardware support





```
#include<stdio.h>
```

```
void main(int argc, char *argv[])  
{
```

```
    int pid;
```

```
    /* Fork another process */  
    pid = fork();
```

```
    if(pid < 0)  
    {
```

```
        //Error occurred  
        fprintf(stderr, "Fork Failed");  
        exit(-1);
```

```
    }  
    else if (pid == 0)
```

```
    {  
        //Child process  
        execlp("/bin/ls", "ls", NULL);
```

```
    }  
    else  
    {
```

```
        //Parent process  
        //Parent will wait for the child to complete  
        wait(NULL);  
        printf("Child complete");  
        exit(0);
```

```
    }  
}
```

GATE Numerical Tip: If fork is called for n times, the number of child processes or new processes created will be: $2^n - 1$.





Process Creation

- ❑ **Parent** process create **children** processes, which, in turn create other processes, forming a tree of processes
- ❑ Generally, process identified and managed via **a process identifier (pid)**
- ❑ Resource sharing
 - ❑ Parent and children share all resources
 - ❑ Children share subset of parent's resources
 - ❑ Parent and child share no resources
- ❑ Execution
 - ❑ Parent and children execute concurrently
 - ❑ Parent waits until children terminate





Process Creation (Cont)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork** system call creates new process
 - **exec** system call used after a **fork** to replace the process' memory space with a new program





Process Operations

-1. Process creation
-fork, spawn

-2. Process Termination


```
main()  
{  
    fork();  
    printf("hello");  
}
```

fork()

-copy of parent, child

-child : 0

-Parent : +ve





Process Operations

-1. Process creation
-fork, spawn

-2. Process Termination

```
main()
```

```
{
```

```
fork();
```

```
fork();
```

```
printf("hello");
```


```
}
```

fork()

-copy of parent, child

-child : 0

-Parent : +ve





Process Operations

-1. Process creation
-fork, spawn

-2. Process Termination


fork()

-copy of parent, child

-child : 0

-Parent : +ve

```
main()
{
fork();
fork();
printf("hello");
}
```




$$2^n - 1$$





Process Creation






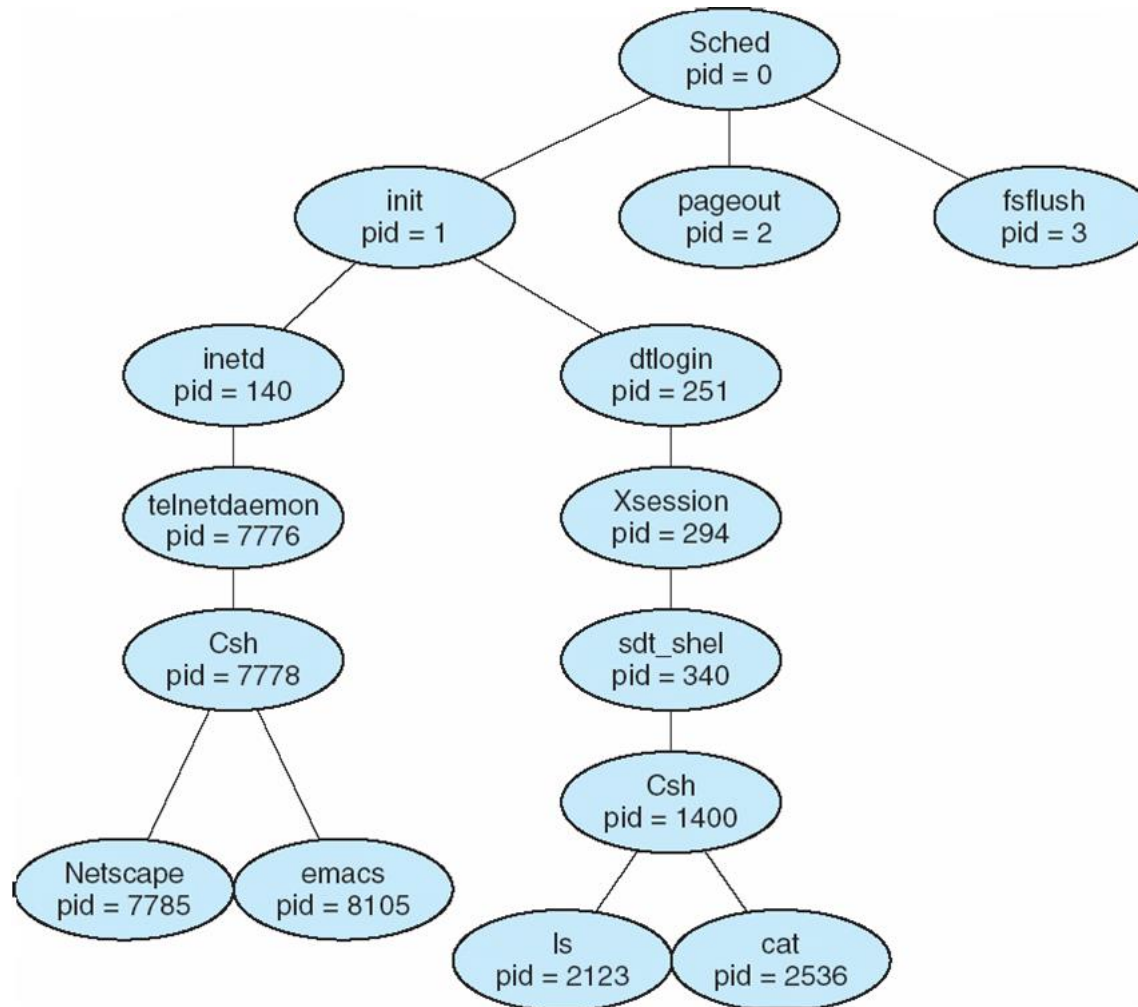
C Program Forking Separate Process

```
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```





A tree of processes on a typical Solaris





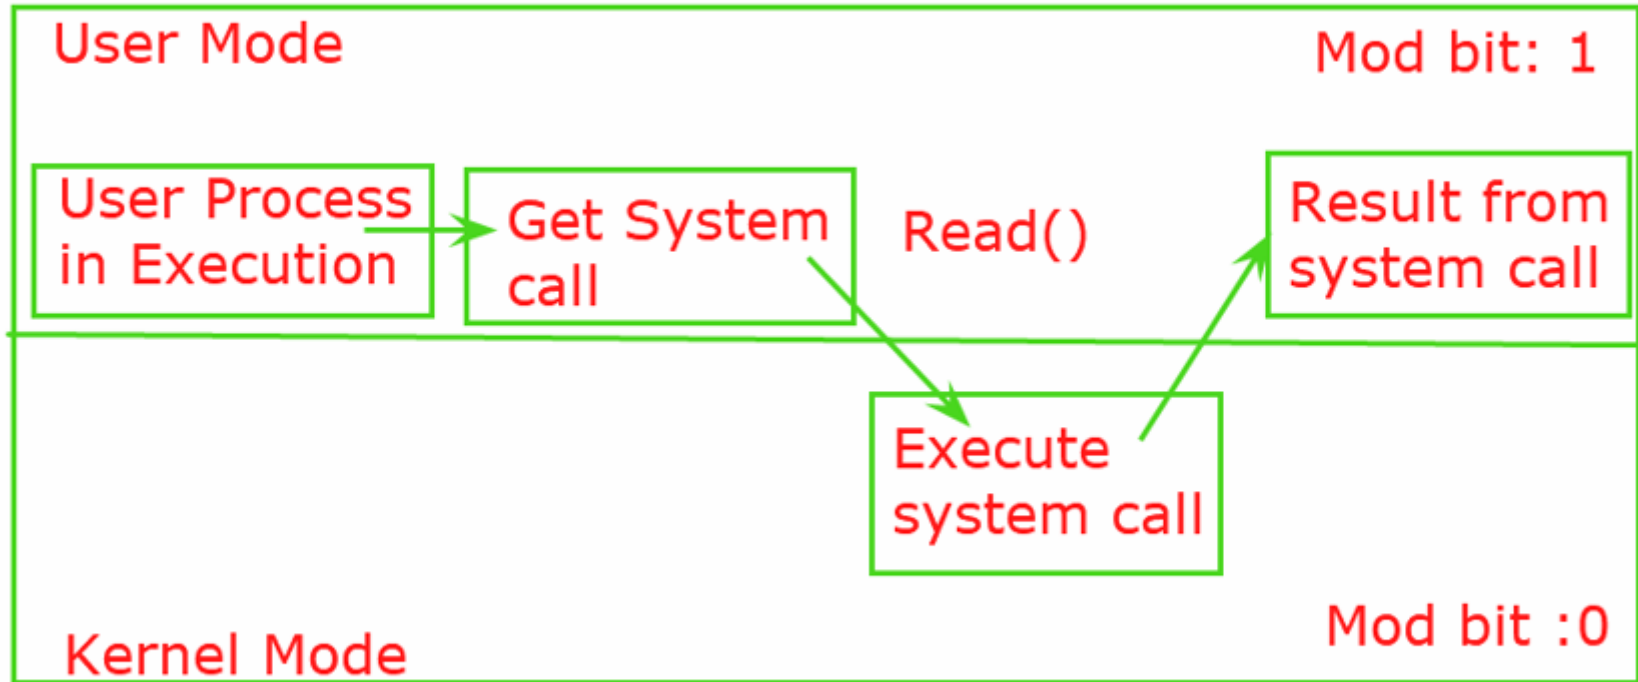
Process Termination

- ❑ Process executes last statement and asks the operating system to delete it (**exit**)
 - ❑ Output data from child to parent (via **wait**)
 - ❑ Process' resources are deallocated by operating system
- ❑ Parent may terminate execution of children processes (**abort**)
 - ❑ Child has exceeded allocated resources
 - ❑ Task assigned to child is no longer required
 - ❑ If parent is exiting
 - ▶ Some operating system do not allow child to continue if its parent terminates
 - All children terminated - **cascading termination**





User Vs Kernel Mode






Interprocess Communication

- ❑ Processes within a system may be **independent** or **cooperating**
- ❑ Cooperating process can affect or be affected by other processes, including sharing data
- ❑ Reasons for cooperating processes:
 - ❑ Information sharing
 - ❑ Computation speedup
 - ❑ Modularity
 - ❑ Convenience
- ❑ Cooperating processes need **interprocess communication (IPC)**
- ❑ Two models of IPC
 - ❑ Shared memory
 - ❑ Message passing






Communications Models



(a)




(b)




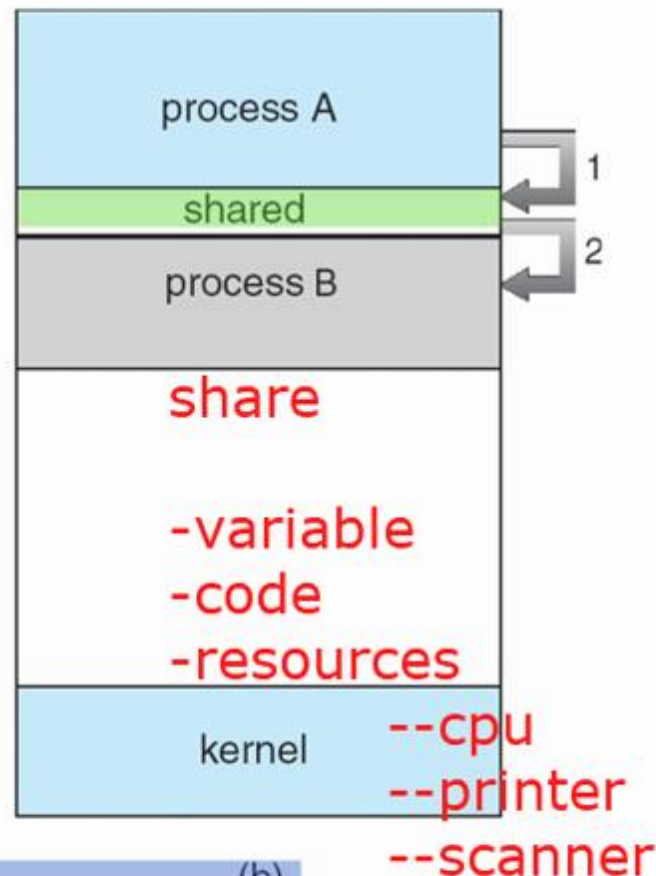


Communications Models

Independent process




Cooperative process



Soln : Synchronization

Ask for same resource--> Racing cond





Cooperating Processes

- ❑ **Independent** process cannot affect or be affected by the execution of another process
- ❑ **Cooperating** process can affect or be affected by the execution of another process
- ❑ Advantages of process cooperation
 - ❑ Information sharing
 - ❑ Computation speed-up
 - ❑ Modularity
 - ❑ Convenience






Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - *unbounded-buffer* places no practical limit on the size of the buffer
 - *bounded-buffer* assumes that there is a fixed buffer size





Producer-Consumer Problem





Producer-Consumer Problem


Consumer

while(ar == n)

20

ar = ar - 1

ar 2 3



Buffer

n = 6, 0-5

Producer

while(ar == n)

10

20

30

ar = ar + 1

40





Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10  
typedef struct {  
    . . .  
} item;  
  
item buffer[BUFFER_SIZE];  
int in = 0;  
int out = 0;
```

- Solution is correct, but can only use BUFFER_SIZE-1 elements





Bounded-Buffer – Producer

```
while (true) {  
    /* Produce an item */  
    while (((in = (in + 1) % BUFFER SIZE count) == out)  
        ; /* do nothing -- no free buffers */  
        buffer[in] = item;  
        in = (in + 1) % BUFFER SIZE;  
    }
```





Bounded Buffer – Consumer

```
while (true) {  
    while (in == out)  
        ; // do nothing -- nothing to consume  
  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
    return item;  
}
```





Interprocess Communication – Message Passing

- ❑ Mechanism for processes to communicate and to synchronize their actions
- ❑ Message system – processes communicate with each other without resorting to shared variables
- ❑ IPC facility provides two operations:
 - ❑ **send**(*message*) – message size fixed or variable
 - ❑ **receive**(*message*)
- ❑ If P and Q wish to communicate, they need to:
 - ❑ establish a *communication link* between them
 - ❑ exchange messages via send/receive
- ❑ Implementation of communication link
 - ❑ physical (e.g., shared memory, hardware bus)
 - ❑ logical (e.g., logical properties)





Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and continue
 - **Non-blocking receive** has the receiver receive a valid message or null






Buffering

- Queue of messages attached to the link; implemented in one of three ways
 1. Zero capacity – 0 messages
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits



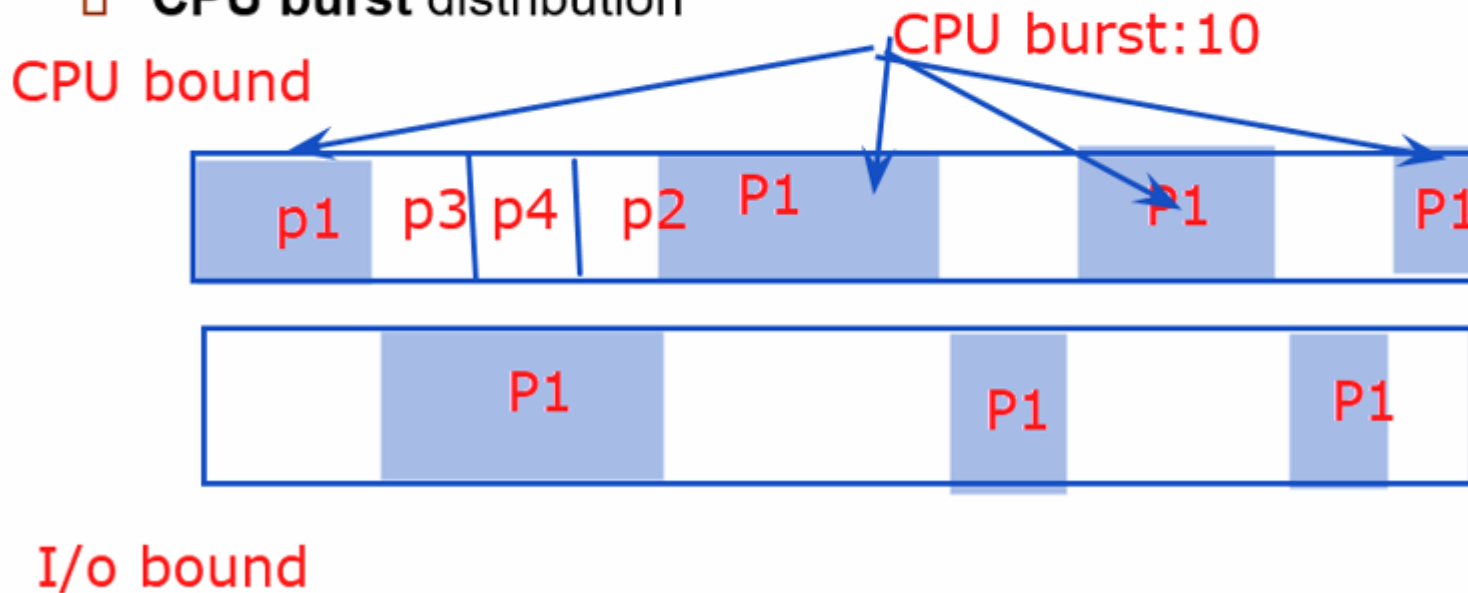
CPU Scheduling






Basic Concepts

- Maximum CPU utilization obtained with multiprogramming
- CPU-I/O Burst Cycle – Process execution consists of a cycle of CPU execution and I/O wait
- **CPU burst** distribution



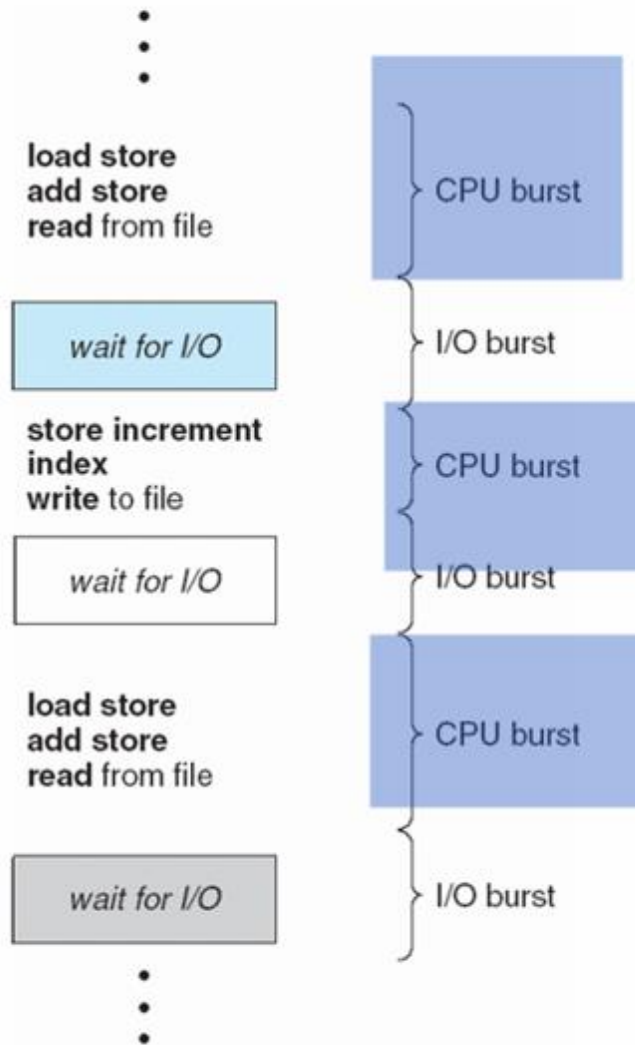


Histogram of CPU-burst Times





Alternating Sequence of CPU And I/O Bursts



1. Running -> waiting
2. Running -> Ready
3. waiting -> Ready

Scheduling





CPU Scheduler

- ❑ Selects from among the **processes in memory that are ready to execute**, and allocates the CPU to one of them
- ❑ CPU scheduling decisions may take place when a process:
 1. Switches from **running to waiting state**
 2. Switches from **running to ready state**
 3. Switches from **waiting to ready**
 4. Terminates
- ❑ Scheduling under 1 and 4 is **nonpreemptive**
- ❑ All other scheduling is **preemptive**





Dispatcher

- ❑ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ❑ switching context
 - ❑ switching to user mode
 - ❑ jumping to the proper location in the user program to restart that program
- ❑ **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running





Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)





Scheduling Algorithm Optimization Criteria

- ❑ Max CPU utilization
- ❑ Max throughput
- ❑ Min turnaround time
- ❑ Min waiting time
- ❑ Min response time

Arrival Time :

Burst Time :

Completion Time:

Turn around Time: {Completion time -arrival time}

Waiting Time : {Turn around - burst time}

Response time : {Process with 1st cpu allotment -arrival time}



Scheduling Algorithm

Pre-emptive

SRTF

LRTF

Round Robin

Priority based

Non- pre-emptive

FCFS

SJF

LJF

HRRN

Multilevel Queue

Priority





First-Come, First-Served (FCFS) Scheduling

Process	Burst Time	comp.T	TAT	WT	Resp T
P₁	24	24	24	0	0
P₂	3	27	27	24	24
P₃	3	30	30	27	27


Sequence: ~~P1~~-~~P2~~-P3



$$\text{AWT: } (0+24+27)/3=17$$

ATAT : 27

Res T : 17





First-Come, First-Served (FCFS) Scheduling

Process	Burst Time
P_1	24
P_2	3
P_3	3

comp.T	TAT	WT	Resp T
24	24	0	0
27	27	24	24
30	30	27	27

Sequence: P2-P3-P1

Arrival time : 0



Gantt Chart

AWT: $(0+24+27)/3=17$

ATAT : 27

Res T : 17





First-Come, First-Served (FCFS) Scheduling

Process	Burst Time	comp.T	TAT	WT	Resp T
P_1	24	24	24	0	0
P_2	3	27	27	24	24
P_3	3	30	30	27	27

Sequence: P2-P3-P1

Arrival time : 0

case 1



$$AWT: (0+24+27)/3=17$$

ATAT : 27

Res T : 17

Convoy Effect

case 2






First-Come, First-Served (FCFS) Scheduling

Process	Burst Time	comp.T	TAT	WT	Resp T
P_1	24	24	24	0	0
P_2	3	27	27	24	24
P_3	3	30	30	27	27

Sequence: P2-P3-P1

Arrival time : 0

case 1



$$\text{AWT: } (0+24+27)/3=17$$


ATAT : 27

Res T : 17

Convoy Effect

$$\text{AWT : } (6+0+3)/3=3$$

case 2





Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- SJF is optimal – gives minimum average waiting time for a given set of processes
 - The difficulty is knowing the length of the next CPU request






Case 2

Process	Burst Time	comp.T	TAT	WT	Resp T
P_1	24	24	24	0	0
P_2	3	27	27	24	24
P_3	3	30	30	27	27

Sequence: P2-P3-P1

Arrival time : 0

case 1



$$AWT: (0+24+27)/3=17$$


ATAT : 27

Res T : 17

Convoy Effect

$$AWT : (6+0+3)/3=3$$

case 2





Example of SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	6
P_2	2.0	8
P_3	4.0	7
P_4	5.0	3





Priority Scheduling


- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- Problem \equiv **Starvation** – low priority processes may never execute
- Solution \equiv **Aging** – as time progresses increase the priority of the process





Priority Scheduling

Process			Burst Time
1	P_1	5	24
2	P_2	0	3
3	P_3	1	3





Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Performance
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high







Example of RR with Time Quantum = 4

Process	Burst Time
P_1	24
P_2	3
P_3	3


Process	Burst Time
1 P_1	24
2 P_2	3
3 P_3	3

Round Robin : time quantum
3 sec





Time Quantum and Context Switch Time





Multilevel Queue

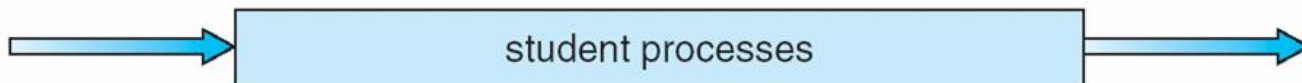


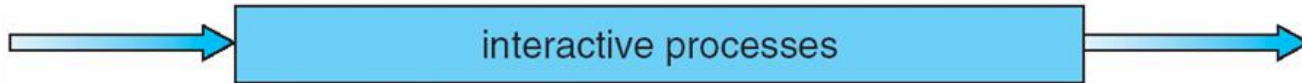
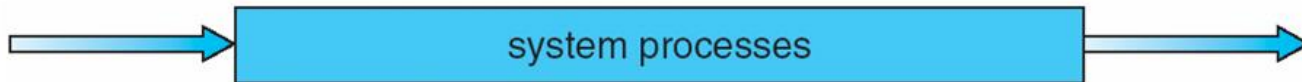
- Ready queue is partitioned into separate queues:
foreground (interactive)
background (batch)
- Each queue has its own scheduling algorithm
 - foreground – RR
 - background – FCFS
- Scheduling must be done between the queues
 - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - 20% to background in FCFS





Multilevel Queue Scheduling

highest priority



lowest priority






Multilevel Feedback Queue

- ❑ A process can move between the various queues; aging can be implemented this way
- ❑ Multilevel-feedback-queue scheduler defined by the following parameters:
 - ❑ number of queues
 - ❑ scheduling algorithms for each queue
 - ❑ method used to determine when to upgrade a process
 - ❑ method used to determine when to demote a process
 - ❑ method used to determine which queue a process will enter when that process needs service





Example of Multilevel Feedback Queue

- Three queues:
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- Scheduling
 - A new job enters queue Q_0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
 - At Q_1 job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .





Multilevel Feedback Queues

